

```
In [0]: %matplotlib inline
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
import nltk
import string
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import model_selection
from sklearn.model_selection import train_test_split
```

```
In [72]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
In [0]: # Getting data into a dataframe
path="/content/drive/My Drive/Colab_Notebooks/ass14/preprocessed_data.csv"
df = pd.read_csv(path)
```

In [49]: `df.head(5)`

Out[49]:

	school_state	teacher_prefix	project_grade_category	teacher_number_of_previously_posted_pro
0	ca	mrs	grades_prek_2	
1	ut	ms	grades_3_5	
2	ca	mrs	grades_prek_2	
3	ga	mrs	grades_prek_2	
4	wa	mrs	grades_3_5	



In [50]: `df.columns`

Out[50]: Index(['school\_state', 'teacher\_prefix', 'project\_grade\_category', 'teacher\_number\_of\_previously\_posted\_projects', 'project\_is\_approved', 'clean\_categories', 'clean\_subcategories', 'essay', 'price'], dtype='object')

In [51]: `df.shape`

Out[51]: (109248, 9)

In [0]: `#df = df.sample(n=40000)  
#project_data=project_data.tail(1000)  
#project_data.shape`

## Splitting dataset

```
In [52]: y=df['project_is_approved']  
y.shape
```

```
Out[52]: (109248,)
```

```
In [53]: features = df.drop(["project_is_approved"],axis=1)  
features.shape
```

```
Out[53]: (109248, 8)
```

```
In [0]: #https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html
```

```
#split the data into train and test fo bag of words
```

```
x_train,x_test,y_train,y_test=model_selection.train_test_split(features,y,test_size=0.33,stratify=y,random_state=0)  
#split train into cross val train and cross val test  
#x_train,x_cv,y_train,y_cv=model_selection.train_test_split(x_t,y_t,test_size=0.3,random_state=0)
```


```
In [55]: print(x_train.shape)  
print("+++++")  
print(x_test.shape)
```

```
(73196, 8)  
+++++  
(36052, 8)
```


```
In [0]: # Preparing Text Data As per Our Model  
x_train["len_essay"] = x_train["essay"].apply(len)  
x_test["len_essay"] = x_test["essay"].apply(len)
```

## Distribution plot of essay dataset

```
In [57]: sns.set()  
ax = sns.distplot(x_train["len_essay"])
```



```
In [59]: ax = sns.distplot(x_test["len_essay"])
```



```
In [0]: from sklearn.feature_extraction.text import CountVectorizer  
from nltk.stem.porter import PorterStemmer  
import re  
import string  
from nltk.corpus import stopwords  
from nltk.stem import PorterStemmer  
from nltk.stem.wordnet import WordNetLemmatizer  
from gensim.models import Word2Vec  
from gensim.models import KeyedVectors  
import pickle
```

```
In [0]: # convert the sentences (strings) into integers
from keras.preprocessing.text import Tokenizer
tokenizer = Tokenizer(num_words=5000)
tokenizer.fit_on_texts(x_train["essay"].tolist())
sequences_train = tokenizer.texts_to_sequences(x_train["essay"])
sequences_test = tokenizer.texts_to_sequences(x_test["essay"])
```

In [61]: sequences\_train

```
Out[61]: [[25,  
6,  
1759,  
30,  
1008,  
16,  
3225,  
12,  
1,  
24,  
802,  
1729,  
3,  
9,  
144,  
19,  
1196,  
1786,  
436,  
645,  
3126,  
46,  
1082,  
762,  
39,  
90,  
12,  
1,  
512,  
4302,  
807,  
191,  
1047,  
108,  
7,  
3786,  
1297,  
277,  
437,  
1,  
333,  
259,  
2707,  
205,  
44,  
36,  
59,  
597,  
568,  
2721,  
6,  
1698,  
129,  
1842,  
146,  
12,  
1,
```

133,  
2505,  
322,  
3231,  
6,  
597,  
9,  
504,  
2118,  
3538,  
2201,  
1351,  
2235,  
1837,  
258,  
597,  
139,  
179,  
333,  
587,  
385,  
66,  
10,  
376,  
2,  
491,  
321,  
1,  
239,  
2235,  
1659,  
78,  
16,  
2,  
24,  
491,  
110,  
631,  
1766,  
76,  
635,  
2,  
950,  
11,  
779,  
1,  
33,  
137,  
826,  
398,  
587,  
4165,  
3181,  
232,  
51,  
547,  
1,



100,  
19,  
3,  
2,  
301,  
34,  
737,  
518,  
61,  
981,  
631,  
2,  
15,  
117,  
980,  
31,  
148,  
117,  
360,  
483,  
9,  
100,  
247,  
936,  
2,  
34,  
83,  
313,  
58,  
2,  
146,  
3,  
556,  
748,  
158,  
2,  
19,  
432,  
2,  
19,  
1,  
199,  
214,  
182,  
45,  
262,  
37,  
58,  
2,  
8,  
723,  
1073,  
86,  
15,  
6,  
54,  
924,

677,  
38,  
270,  
232,  
587,  
516,  
5,  
85,  
2,  
15,  
2806,  
942,  
583,  
1,  
8,  
40,  
4521,  
2011,  
1572,  
2,  
15,  
879,  
929,  
650,  
80,  
531,  
264,  
2,  
15,  
1056,  
2,  
603,  
173,  
453,  
1053,  
536,  
40,  
1,  
375,  
2,  
15,  
2868,  
825,  
117,  
1,  
360,  
1659,  
432,  
66,  
766,  
398,  
89,  
2589,  
2,  
16,  
50,  
131,

```
3,  
1792,  
2360,  
1011,  
40,  
63,  
54,  
182,  
113,  
33,  
411,  
385,  
182,  
113,  
33,  
385,  
6,  
13],  
...]
```

```
In [62]: print("No. of datapoints in X_train :",len(x_train))  
print("No. of datapoints in X_test :",len(x_test))  
print("Shape of Y_train :",y_train.shape)  
print("Shape of Y_test :",y_test.shape)
```

```
No. of datapoints in X_train : 73196  
No. of datapoints in X_test : 36052  
Shape of Y_train : (73196,)  
Shape of Y_test : (36052,)
```

```
In [63]: # get word -> integer mapping  
word2idx = tokenizer.word_index  
print('Found %s unique tokens.' % len(word2idx))
```

```
Found 48499 unique tokens.
```

```
In [0]: # importing required libraries
import warnings
warnings.filterwarnings("ignore")
import pandas as pd
import numpy as np
from keras.layers import Input, Embedding, LSTM, Dropout, BatchNormalization,
Dense, concatenate, Flatten, Conv1D, MaxPool1D, LeakyReLU, ELU, SpatialDropout
1D, MaxPooling1D, GlobalAveragePooling1D, GlobalMaxPooling1D
from keras.preprocessing.text import Tokenizer, one_hot
from keras.preprocessing.sequence import pad_sequences
from keras.models import Model, load_model
from keras import regularizers
from keras.optimizers import *
from keras.callbacks import ModelCheckpoint, EarlyStopping, TensorBoard, Reduc
eLROnPlateau
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.metrics import roc_auc_score
import tensorflow as tf
import matplotlib.pyplot as plt
%matplotlib inline
import re
from tqdm import tqdm
from sklearn.preprocessing import LabelEncoder
import seaborn as sns
import pickle
```

```
In [65]: # truncate and/or pad input sequences
max_review_length = 800
encoded_train = pad_sequences(sequences_train,maxlen=max_review_length,padding
='post', truncating='post')
encoded_test = pad_sequences(sequences_test, maxlen=max_review_length,padding=
'post', truncating='post')
print('Shape of train data tensor:', encoded_train.shape)
print('Shape of test data tensor:', encoded_test.shape)

print(encoded_train[1])
```

```
Shape of train data tensor: (73196, 800)
```

```
Shape of test data tensor: (36052, 800)
```

[illegible]

```

0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0]

```

```

In [0]: # Loading Embedding File
pickle_in = open("glove_vectors", "rb")
glove_words = pickle.load(pickle_in)

```

```

In [77]: MAX_VOCAB_SIZE=5000
num_words = min(MAX_VOCAB_SIZE, len(word2idx) + 1)
embedding_matrix = np.zeros((num_words, 300))
for word, i in word2idx.items():
    if i < MAX_VOCAB_SIZE:
        embedding_vector = glove_words.get(word)
        if embedding_vector is not None:
            # words not found in embedding index will be all zeros.
            embedding_matrix[i] = embedding_vector

print(num_words)
print("+++++")
print(embedding_matrix.shape)

```

```

5000
+++++++
(5000, 300)

```

```
In [78]: # Load pre-trained word embeddings into an Embedding layer
# note that we set trainable = False so as to keep the embeddings fixed
MAX_SEQUENCE_LENGTH=800
embedding_layer = Embedding(
    num_words,
    300,
    weights=[embedding_matrix],
    input_length=MAX_SEQUENCE_LENGTH,
    trainable=False
)
input_text = Input(shape=(MAX_SEQUENCE_LENGTH,), name="input_text")
x = embedding_layer(input_text)
x = LSTM(128, recurrent_dropout=0.5, kernel_regularizer=regularizers.l2(0.001),
return_sequences=True)(x) # dropout=0.5
# x = SpatialDropout1D(0.5)(x)
flatten_1 = Flatten()(x)
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:66: The name tf.get\_default\_graph is deprecated. Please use tf.compat.v1.get\_default\_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:4432: The name tf.random\_uniform is deprecated. Please use tf.random.uniform instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:190: The name tf.get\_default\_session is deprecated. Please use tf.compat.v1.get\_default\_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow\_backend.py:3733: calling dropout (from tensorflow.python.ops.nn\_ops) with keep\_prob is deprecated and will be removed in a future version.  
Instructions for updating:  
Please use `rate` instead of `keep\_prob`. Rate should be set to `rate = 1 - keep\_prob`.

## categorical variable



```
In [79]: # Now will prepare all the remaining categorical features
# Teacher Prefix
no_of_unique_prefix = x_train["teacher_prefix"].nunique()
embedding_size_prefix = int(min(np.ceil((no_of_unique_prefix)/2), 50 ))
print('Unique Categories:', no_of_unique_prefix, 'Embedding Size:', embedding_size_prefix)

# Defining Input and Embedding Layer for the same

input_prefix = Input(shape=(1,),name="teacher_prefix")
embedding_prefix = Embedding(no_of_unique_prefix,embedding_size_prefix,name="embedding_pre",trainable=True)(input_prefix)
flatten_2 = Flatten()(embedding_prefix)

lb = LabelEncoder()
encoder_prefix_train = lb.fit_transform(x_train["teacher_prefix"])
# encoder_prefix_cv = lb.transform(X_cv["teacher_prefix"])
encoder_prefix_test = lb.transform(x_test["teacher_prefix"])
```

Unique Categories: 5 Embedding Size: 3

```
In [80]: # School State
no_of_unique_state = x_train["school_state"].nunique()
embedding_size_state= int(min(np.ceil((no_of_unique_state)/2), 50 ))
print('Unique Categories:', no_of_unique_state, 'Embedding Size:', embedding_size_state)

# Defining Input and Embedding Layer for the same

input_state = Input(shape=(1,),name="school_prefix")
embedding_state = Embedding(no_of_unique_state,embedding_size_state,name="embedding_state",trainable=True)(input_state)
flatten_3 = Flatten()(embedding_state)

encoder_state_train = lb.fit_transform(x_train["school_state"])
# encoder_state_cv = lb.transform(X_cv["school_state"])
encoder_state_test = lb.transform(x_test["school_state"])
```

Unique Categories: 51 Embedding Size: 26

```
In [81]: # For project_grade_category
no_of_unique_grade = x_train["project_grade_category"].nunique()
embedding_size_grade = int(min(np.ceil((no_of_unique_grade)/2), 50 ))
print('Unique Categories:', no_of_unique_grade, 'Embedding Size:', embedding_size_grade)

# Defining Input and Embedding Layer for the same

input_grade= Input(shape=(1,),name="grade_cat")
embedding_grade = Embedding(no_of_unique_grade,embedding_size_grade,name="emb_grade",trainable=True)(input_grade)
flatten_4 = Flatten()(embedding_grade)

encoder_grade_train = lb.fit_transform(x_train["project_grade_category"])
# encoder_grade_cv = lb.transform(X_cv["project_grade_category"])
encoder_grade_test = lb.transform(x_test["project_grade_category"])
```

Unique Categories: 4 Embedding Size: 2

```
In [82]: # For clean_categories
no_of_unique_subcat = x_train["clean_categories"].nunique()
embedding_size_subcat = int(min(np.ceil((no_of_unique_subcat)/2), 50 ))
print('Unique Categories:', no_of_unique_subcat, 'Embedding Size:', embedding_size_subcat)

# Defining Input and Embedding Layer for the same

input_subcat= Input(shape=(1,),name="sub_cat")
embedding_subcat = Embedding(no_of_unique_subcat,embedding_size_subcat,name="emb_subcat",trainable=True)(input_subcat)
flatten_5 = Flatten()(embedding_subcat)

# encoder_subcat_train = lb.fit_transform(x_train["clean_categories"])
# encoder_subcat_cv = lb.transform(X_cv["clean_categories"])
# encoder_subcat_test = lb.transform(x_test["clean_categories"])
le = LabelEncoder()
le.fit(x_train["clean_categories"])
x_test["clean_categories"] = x_test["clean_categories"].map(lambda s: '<unknown>' if s not in le.classes_ else s)
# X_cv["clean_categories"] = X_cv["clean_categories"].map(lambda s: '<unknown>' if s not in le.classes_ else s)
le.classes_ = np.append(le.classes_, '<unknown>')
encoder_subcat_train = le.transform(x_train["clean_categories"])
encoder_subcat_test= le.transform(x_test["clean_categories"])
# encoder_subcat_cv = le.transform(X_cv["clean_categories"])
```

Unique Categories: 51 Embedding Size: 26

```
In [83]: # For clean_subcategories
no_of_unique_subcat_1 = x_train["clean_subcategories"].nunique()
embedding_size_subcat_1 = int(min(np.ceil((no_of_unique_subcat_1)/2), 50 ))
print('Unique Categories:', no_of_unique_subcat_1, 'Embedding Size:', embedding_size_subcat_1)

# Defining Input and Embedding Layer for the same

input_subcat_1= Input(shape=(1,),name="sub_cat_1")
embedding_subcat_1 = Embedding(no_of_unique_subcat_1+1,embedding_size_subcat_1,
name="emb_subcat_1",trainable=True)(input_subcat_1)#adding +1
flatten_6 = Flatten()(embedding_subcat_1)

le = LabelEncoder()
le.fit(x_train["clean_subcategories"])
x_test["clean_subcategories"] = x_test["clean_subcategories"].map(lambda s: '<unknown>' if s not in le.classes_ else s)
# X_cv["clean_subcategories"] = X_cv["clean_subcategories"].map(lambda s: '<unknown>' if s not in le.classes_ else s)
le.classes_ = np.append(le.classes_, '<unknown>')
encoder_subcat_1_train = le.transform(x_train["clean_subcategories"])
encoder_subcat_1_test= le.transform(x_test["clean_subcategories"])
# encoder_subcat_1_cv = le.transform(X_cv["clean_subcategories"])
```

Unique Categories: 390 Embedding Size: 50

## numerical data

```
In [0]: # Now we will prepare numerical features for our model
num_train_1=x_train['len_essay'].values.reshape(-1, 1)
num_train_2=x_train['price'].values.reshape(-1, 1)
num_train_3=x_train['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1)

num_test_1=x_test['len_essay'].values.reshape(-1, 1)
num_test_2=x_test['price'].values.reshape(-1, 1)
num_test_3=x_test['teacher_number_of_previously_posted_projects'].values.reshape(-1, 1)

num_train=np.concatenate((num_train_1,num_train_2,num_train_3),axis=1)

num_test=np.concatenate((num_test_1,num_test_2,num_test_3),axis=1)
```

```
In [0]: from sklearn.preprocessing import StandardScaler
norm=StandardScaler()
norm_train=norm.fit_transform(num_train)
norm_test=norm.transform(num_test)
```

```
In [0]: # Defining the Input and Embedding Layer for the same
num_feats = Input(shape=(3,),name="numerical_features")
num_feats_ = Dense(100,activation="relu",kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001))(num_feats)
```

```
In [0]: x_concatenate = concatenate([flatten_1,flatten_2,flatten_3,flatten_4,flatten_5,flatten_6,num_feats_])
```

```
In [88]: print("Building Model-1")

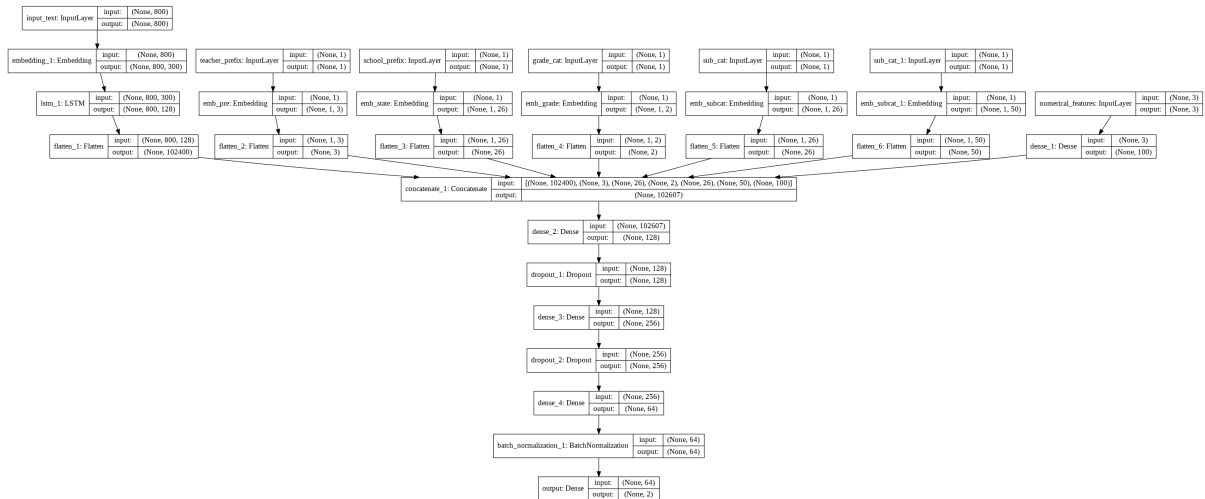
# x_concatenate = BatchNormalization()(x_concatenate)
x = Dense(128,activation="relu", kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001))(x_concatenate)
# x=LeakyReLU(alpha=0.3)(x)
x=Dropout(0.5)(x)
x = Dense(256,activation="relu",kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001))(x)
# x=LeakyReLU(alpha=0.3)(x)
x=Dropout(0.5)(x)
x = Dense(64,activation="relu", kernel_initializer="he_normal",kernel_regularizer=regularizers.l2(0.001))(x)
x = BatchNormalization()(x)
# x=LeakyReLU(alpha=0.3)(x)
output = Dense(2, activation='softmax', name='output')(x)
model_1 = Model(inputs=[input_text,input_prefix,input_state,input_grade,
                        input_subcat,input_subcat_1,num_feats],outputs=[output
])
```

Building Model-1

```
In [89]: # https://github.com/mmortazavi/EntityEmbedding-Working_Example/blob/master/EntityEmbedding.ipynb
#https://stackoverflow.com/questions/36886711/keras-runtimeerror-failed-to-import-pydot-after-installing-graphviz-and-pydot
from keras.utils import plot_model
import keras
import pydotplus
from keras.utils.vis_utils import model_to_dot
#keras.utils.vis_utils.pydot = pydot

#import pydot_ng as pydot
plot_model(model_1, show_shapes=True, show_layer_names=True, to_file='model_1.png')
from IPython.display import Image
Image(retina=True, filename='model_1.png')
```

Out[89]:



```
In [0]: train_data_1 = [encoded_train,encoder_prefix_train,encoder_state_train,
                        encoder_grade_train,encoder_subcat_train,encoder_subcat_1_train,
                        norm_train]
test_data_1 = [encoded_test,encoder_prefix_test,encoder_state_test,encoder_grade_test,
               encoder_subcat_test,encoder_subcat_1_test,norm_test]

from keras.utils import np_utils
Y_train = np_utils.to_categorical(y_train, 2)
Y_test = np_utils.to_categorical(y_test, 2)
```

```
In [0]: checkpoint_1 = ModelCheckpoint("model_1.h5",
                                       monitor="val_loss",
                                       mode="min",
                                       save_best_only = True,
                                       verbose=1)

earlystop_1 = EarlyStopping(monitor = 'val_loss',
                             mode="min",
                             min_delta = 0,
                             patience = 2,
                             verbose = 1,
                             restore_best_weights = True)

reduce_lr_1 = ReduceLROnPlateau(monitor = 'val_loss', factor = 0.2, patience =
1, verbose = 1, min_delta = 0.0001)

tensorboard_1 = TensorBoard(log_dir='graph_1', histogram_freq=0, batch_size=51
2, write_graph=True, write_grads=False, write_images=False, embeddings_freq=0,
embeddings_layer_names=None, embeddings_metadata=None, embeddings_data=None, u
pdate_freq='epoch')

callbacks_1 = [checkpoint_1,earlystop_1,tensorboard_1,reduce_lr_1]
```

```
In [0]: # Defining Custom ROC-AUC Metrics
from sklearn.metrics import roc_auc_score

def auc1(y_true, y_pred):
    if len(np.unique(y_true[:,1])) == 1:
        return 0.5
    else:
        return roc_auc_score(y_true, y_pred)

def auROC(y_true, y_pred):
    return tf.py_func(auc1, (y_true, y_pred), tf.double)
```

```
In [0]: adam = Adam(lr=0.001, beta_1=0.9, beta_2=0.999, epsilon=None, decay=0.0, amsgr
ad=False)
```

```
In [94]: model_1.compile(optimizer=adam, loss='categorical_crossentropy', metrics=[auroc])
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

```
WARNING:tensorflow:From <ipython-input-92-a7e6cba44e56>:10: py_func (from tensorflow.python.ops.script_ops) is deprecated and will be removed in a future version.
```

Instructions for updating:

tf.py\_func is deprecated in TF V2. Instead, there are two options available in V2.

- tf.py\_function takes a python function which manipulates tf eager tensors instead of numpy arrays. It's easy to convert a tf eager tensor to

o

- an ndarray (just call `tensor.numpy()`) but having access to eager tensors means `tf.py_function`'s` can use accelerators such as GPUs as well as being differentiable using a gradient tape.

- `tf.numpy_function` maintains the semantics of the deprecated `tf.py_func` (it is not differentiable, and manipulates numpy arrays). It drops the stateful argument making all functions stateful.

```
In [54]: history_1 = model_1.fit(train_data_1,Y_train,batch_size=512,  
                                epochs=20,validation_data=(test_data_1,Y_test),callbac  
                                ks=callbacks_1)
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/math\_grad.py:1250: add\_dispatch\_support.<locals>.wrapper (from tensorflow.python.ops.array\_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 73196 samples, validate on 36052 samples

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1122: The name tf.summary.merge\_all is deprecated. Please use tf.compat.v1.summary.merge\_all instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/callbacks.py:1125: The name tf.summary.FileWriter is deprecated. Please use tf.compat.v1.summary.FileWriter instead.

Epoch 1/20

73196/73196 [=====] - 310s 4ms/step - loss: 1.3721 - auroc: 0.5209 - val\_loss: 0.9704 - val\_auroc: 0.6374

Epoch 00001: val\_loss improved from inf to 0.97041, saving model to model\_1.h5

Epoch 2/20

73196/73196 [=====] - 308s 4ms/step - loss: 0.7985 - auroc: 0.6546 - val\_loss: 0.7230 - val\_auroc: 0.7188

Epoch 00002: val\_loss improved from 0.97041 to 0.72301, saving model to model\_1.h5

Epoch 3/20

73196/73196 [=====] - 307s 4ms/step - loss: 0.6342 - auroc: 0.7148 - val\_loss: 0.5958 - val\_auroc: 0.7400

Epoch 00003: val\_loss improved from 0.72301 to 0.59585, saving model to model\_1.h5

Epoch 4/20

73196/73196 [=====] - 307s 4ms/step - loss: 0.5477 - auroc: 0.7338 - val\_loss: 0.5348 - val\_auroc: 0.7468

Epoch 00004: val\_loss improved from 0.59585 to 0.53484, saving model to model\_1.h5

Epoch 5/20

73196/73196 [=====] - 307s 4ms/step - loss: 0.4991 - auroc: 0.7397 - val\_loss: 0.4841 - val\_auroc: 0.7482

Epoch 00005: val\_loss improved from 0.53484 to 0.48410, saving model to model\_1.h5

Epoch 6/20

73196/73196 [=====] - 307s 4ms/step - loss: 0.4659 - auroc: 0.7465 - val\_loss: 0.4550 - val\_auroc: 0.7512

Epoch 00006: val\_loss improved from 0.48410 to 0.45502, saving model to model\_1.h5

Epoch 7/20

73196/73196 [=====] - 305s 4ms/step - loss: 0.4446 - auroc: 0.7497 - val\_loss: 0.4410 - val\_auroc: 0.7508

Epoch 00007: val\_loss improved from 0.45502 to 0.44101, saving model to model\_1.h5

Epoch 8/20

73196/73196 [=====] - 305s 4ms/step - loss: 0.4304 -  
auroc: 0.7486 - val\_loss: 0.4317 - val\_auroc: 0.7484

Epoch 00008: val\_loss improved from 0.44101 to 0.43174, saving model to model\_1.h5

Epoch 9/20

73196/73196 [=====] - 305s 4ms/step - loss: 0.4224 -  
auroc: 0.7507 - val\_loss: 0.4155 - val\_auroc: 0.7521

Epoch 00009: val\_loss improved from 0.43174 to 0.41553, saving model to model\_1.h5

Epoch 10/20

73196/73196 [=====] - 305s 4ms/step - loss: 0.4146 -  
auroc: 0.7540 - val\_loss: 0.4182 - val\_auroc: 0.7510

Epoch 00010: val\_loss did not improve from 0.41553

Epoch 00010: ReduceLROnPlateau reducing learning rate to 0.00020000000949949026.

Epoch 11/20

73196/73196 [=====] - 304s 4ms/step - loss: 0.3966 -  
auroc: 0.7656 - val\_loss: 0.4018 - val\_auroc: 0.7538

Epoch 00011: val\_loss improved from 0.41553 to 0.40177, saving model to model\_1.h5

Epoch 12/20

73196/73196 [=====] - 304s 4ms/step - loss: 0.3879 -  
auroc: 0.7720 - val\_loss: 0.4001 - val\_auroc: 0.7539

Epoch 00012: val\_loss improved from 0.40177 to 0.40010, saving model to model\_1.h5

Epoch 13/20

73196/73196 [=====] - 304s 4ms/step - loss: 0.3861 -  
auroc: 0.7751 - val\_loss: 0.4014 - val\_auroc: 0.7520

Epoch 00013: val\_loss did not improve from 0.40010

Epoch 00013: ReduceLROnPlateau reducing learning rate to 4.0000001899898055e-05.

Epoch 14/20

73196/73196 [=====] - 303s 4ms/step - loss: 0.3778 -  
auroc: 0.7859 - val\_loss: 0.4003 - val\_auroc: 0.7516

Epoch 00014: val\_loss did not improve from 0.40010

Restoring model weights from the end of the best epoch

Epoch 00014: ReduceLROnPlateau reducing learning rate to 8.000000525498762e-06.

Epoch 00014: early stopping

```
In [55]: print(model_1.summary())
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_text (InputLayer)	(None, 800)	0	
embedding_1 (Embedding) [0][0]	(None, 800, 300)	1500000	input_text
teacher_prefix (InputLayer)	(None, 1)	0	
school_prefix (InputLayer)	(None, 1)	0	
grade_cat (InputLayer)	(None, 1)	0	
sub_cat (InputLayer)	(None, 1)	0	
sub_cat_1 (InputLayer)	(None, 1)	0	
lstm_1 (LSTM) [0][0]	(None, 800, 128)	219648	embedding_1
emb_pre (Embedding) ix[0][0]	(None, 1, 3)	15	teacher_pref
emb_state (Embedding) x[0][0]	(None, 1, 26)	1326	school_prefi
emb_grade (Embedding) [0]	(None, 1, 2)	8	grade_cat[0]
emb_subcat (Embedding) [0]	(None, 1, 26)	1326	sub_cat[0]
emb_subcat_1 (Embedding) [0]	(None, 1, 50)	19550	sub_cat_1[0]
numerical_features (InputLayer)	(None, 3)	0	
flatten_1 (Flatten)	(None, 102400)	0	lstm_1[0][0]

flatten_2 (Flatten) [0]	(None, 3)	0	emb_pre[0]
flatten_3 (Flatten) [0]	(None, 26)	0	emb_state[0]
flatten_4 (Flatten) [0]	(None, 2)	0	emb_grade[0]
flatten_5 (Flatten) [0][0]	(None, 26)	0	emb_subcat
flatten_6 (Flatten) [0][0]	(None, 50)	0	emb_subcat_1
dense_1 (Dense) atures[0][0]	(None, 100)	400	numerical_fe
concatenate_1 (Concatenate) [0]	(None, 102607)	0	flatten_1[0]
[0]			flatten_2[0]
[0]			flatten_3[0]
[0]			flatten_4[0]
[0]			flatten_5[0]
[0]			flatten_6[0]
[0]			dense_1[0]
[0]			
dense_2 (Dense) 1[0][0]	(None, 128)	13133824	concatenate_
dropout_1 (Dropout) [0]	(None, 128)	0	dense_2[0]
dense_3 (Dense) [0]	(None, 256)	33024	dropout_1[0]
dropout_2 (Dropout) [0]	(None, 256)	0	dense_3[0]

dense_4 (Dense)	(None, 64)	16448	dropout_2[0]
<hr/>			
batch_normalization_1 (BatchNor	(None, 64)	256	dense_4[0]
[0]			
<hr/>			
output (Dense)	(None, 2)	130	batch_normal
ization_1[0][0]			
=====			
=====			
Total params: 14,925,955			
Trainable params: 13,425,827			
Non-trainable params: 1,500,128			
<hr/>			
None			

```
In [0]: my_model = load_model("model_1.h5", custom_objects={"auroc": auroc})
```

```
In [0]: project_status = {0: "Rejected", 1: "Approved"}
```

```
In [0]: Y_pred = my_model.predict(test_data_1, batch_size=512)
```

```
In [0]: # took the function from https://nbviewer.jupyter.org/github/pranaya-mathur/Human-Activity-Recognition/blob/master/Human_Activity_Recognition.ipynb
def confusion_matrix(Y_true, Y_pred):
    Y_true = pd.Series([project_status[y] for y in np.argmax(Y_test, axis=1)])
    Y_pred = pd.Series([project_status[y] for y in np.argmax(Y_pred, axis=1)])

    return pd.crosstab(Y_true, Y_pred, rownames=['True'], colnames=['Pred'])
```

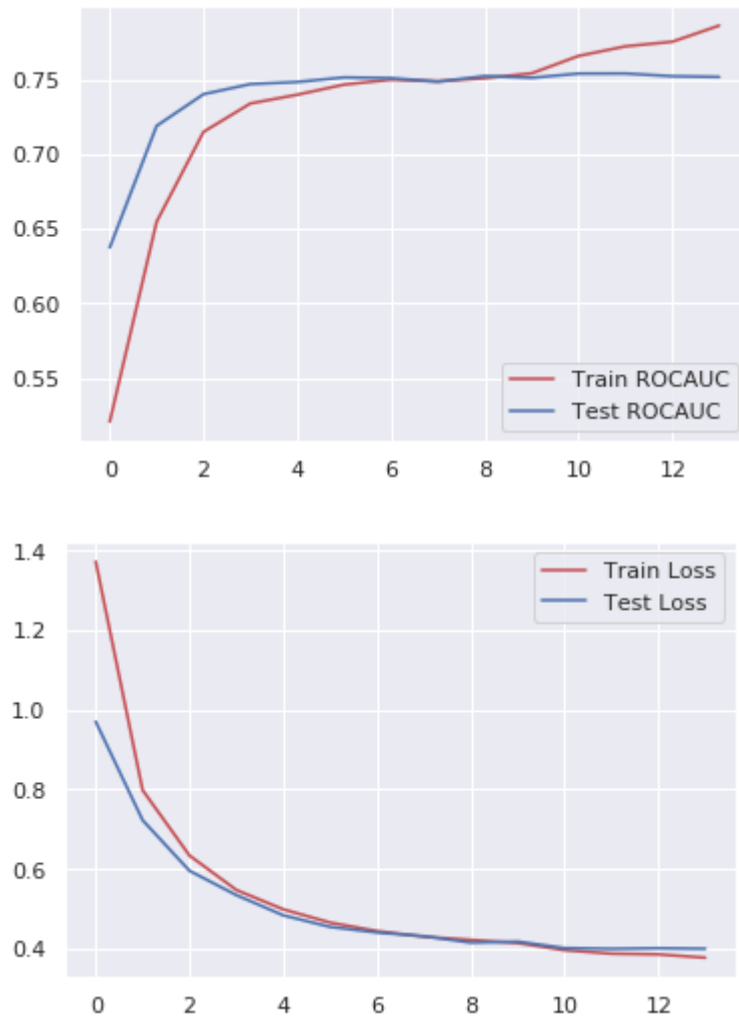
```
In [60]: results = confusion_matrix(Y_test, Y_pred)
results
```

Out[60]:

	Pred Approved	Pred Rejected
True		
Approved	30162	431
Rejected	4825	634

```
In [61]: plt.plot(history_1.history['auroc'], 'r')
plt.plot(history_1.history['val_auroc'], 'b')
plt.legend({'Train ROCAUC': 'r', 'Test ROCAUC': 'b'})
plt.show()

plt.plot(history_1.history['loss'], 'r')
plt.plot(history_1.history['val_loss'], 'b')
plt.legend({'Train Loss': 'r', 'Test Loss': 'b'})
plt.show()
```



Model is overfitting because at end, train AUC(0.786) is higher than test AUC(0.76)

In [0]: